

System Design Interview Framework: 'Build X'

0. Real-World Framing (1 min)

- **Purpose:** Acknowledge the gap between interview exercises and real-world system design.
- **Key Point:** In practice, you design systems with the team you have. Familiarity, existing tooling, and organizational risk tolerance matter more than textbook architectures.
- **Suggested Opening Statement**

"This interview is designed to evaluate my fluency in system design, familiarity with common technologies, and ability to weigh trade-offs. But in the real world, architecture choices are grounded in what your team knows and has the capacity to maintain. So while I'll explore leading practices and trade-offs, I'll stay grounded in pragmatic choices a team could realistically execute."

1. Scope & Assumptions (5 min)


- **1.1 Business Motivation**
 - What problem are we solving, and why now?
 - Strategic objectives and key stakeholders.
 - Organizational success criteria (e.g. revenue uplift, cost reduction, user growth).
- **1.2 Users & Use Cases**
 - Primary personas and their goals.
 - Core journeys (MVP vs. full feature set).
 - Platform constraints (web, iOS, Android, public API).
- **1.3 Budget & Monetization**
 - Phase-appropriate budget limits (MVP vs. scale).
 - Revenue models: ads, subscriptions, transactions.
 - Unit-economics impact of design choices (cost per user/transaction).
- **1.4 Design & Optimization Goals**
 - Performance targets (P50/P95 latency, throughput).
 - Reliability commitments (SLAs, [error budgets](#)).
 - Scalability thresholds (peak users/day, requests/sec).
- **1.5 Success & Failure Metrics**
 - Technical KPIs: latency percentiles, availability, error rate.
 - Business KPIs: conversion, engagement, retention.
 - Operational health: [MTTR](#), incident count.
 - Cost efficiency: infrastructure cost per user/transaction.

- 1.6 Other Considerations

- Internationalization & localization.
- Data privacy & regulatory compliance ([GDPR](#), [HIPAA](#)).
- Supported versions/platforms and upgrade strategy.

2. High-Level Architecture (3 min)

Before diving into specific components, it's helpful to align on the high-level architecture and how core traffic flows through the system.



- 2.1 Tiered Architecture (Four-Tier Model)

Client → Edge → Business → Data

- **Client Tier**

- Mobile/Web clients handle UI rendering, local caching, and request shaping.

- **Edge Tier**

- Ingress: [CDN](#), [TLS termination](#), [L7 routing](#), basic caching.
- API Gateway (optional): Authentication, rate-limits, quotas, request shaping.

- **Business Tier**

- BFFs (Backends for Frontends): Client-specific adapters, payload aggregation.
- Core Services: Domain logic, orchestration, long-running jobs.
- Service Mesh: [mTLS](#), retries, circuit-breaking, observability.

- **Data Tier**

- Databases, object storage, queues, and external APIs.

- 2.2 Request Flow

1. **Client** issues a request →
2. **Ingress** handles TLS, routing & CDN cache →

3. (Public API only) **API Gateway** enforces auth & quotas →
 4. **BFF** transforms and aggregates →
 5. **Core Services** run business logic (sync or via events) →
 6. **Data Tier** persists/retrieves state →
 7. Response returns via the same path.
-

3. Frontend (3min)

- 3.1 Platforms

- Web: [React](#), [Vue](#), [Angular](#), [Next.js](#).
- Native: [Swift](#) (iOS), [Kotlin](#)/Java (Android).
- Cross-platform: React Native, [Flutter](#), [.NET MAUI](#).

- 3.2 Data Fetching

- [REST](#) vs. [GraphQL](#) based on payload flexibility.
- Use REST for external/public APIs and [gRPC](#) for internal service-to-service calls that benefit from type safety and lower overhead.
- Polling for semi-real-time updates.
- WebSockets or [SSE](#) for persistent two-way communication.

- 3.3 API Design Principles

- Use REST for straightforward resource-based APIs and widespread tooling compatibility.
- Use GraphQL for complex data-fetching needs where clients benefit from flexible payload shaping and bundling.
- Follow consistent versioning strategy (e.g., URI versioning or headers).
- Define pagination, filtering, and sorting behaviors explicitly.
- Treat APIs as products: document clearly, support deprecation, and consider developer experience (DX) across web, mobile, and partner clients.

- 3.4 Client-Side Caching & Offline Support

- In-memory storage and local persistence ([IndexedDB](#) for web; [AsyncStorage](#) or [SQLite](#) for mobile/main apps).
- Offline queue for writes and background sync.
- Cache invalidation strategies ([stale-while-revalidate](#), TTL).

- 3.5 Session Management & Notifications

- Token storage & refresh patterns (cookies, localStorage).
- Secure handling of refresh tokens (HttpOnly cookies, secure storage).
- Push notifications setup ([APNs](#), [FCM](#)) and in-app notification handlers.

- 3.6 Ownership

- Typically owned by frontend or full-stack engineering squads.
 - Collaborates closely with UX, product, and backend/BFF teams.
-

4. Edge / Gateway Layer (3 min)

- **4.1 Purpose and Role:** This layer sits between your clients and your backend systems, providing security, traffic management, and protocol translation so frontend and backend teams can iterate independently.
 - **4.2 Edge Components**
 - **Ingress Proxy**
 - CDN & Global Load Balancer (e.g. [Fastly](#), [CloudFront](#)).
 - TLS termination, health checks, basic L7 routing.
 - **API Gateway** (*optional for public APIs*)
 - Authentication & authorization ([JWT/OAuth introspection](#), API keys).
 - Rate limiting, IP allow/deny lists, request validation (schema, size).
 - Request shaping (body size quotas, header transformations) .
 - **Cross-Cutting Edge Services**
 - Feature-flag evaluation ([LaunchDarkly](#), [Flagsmith](#)).
 - Token bootstrapping for service-mesh sidecars.
 - Edge compute functions ([Lambda@Edge](#), [Cloudflare Workers](#)).
 - **4.3 Core Responsibilities**
 - **Security and Compliance**
 - Enforce [authN/authZ](#) before any backend logic runs.
 - Protect against [DDoS](#), [OWASP Top 10](#), and [IP-based threats](#) like IP spoofing.
 - **Traffic Management**
 - Rate limits, quotas, circuit-breaking at the edge.
 - Canary or blue/green routing for safe deployments.
See Section 8.2 ("Deployment Strategies") for additional detail.
 - **Request Routing & Protocol Translation**
 - Route to the correct backend service or BFF.
 - Translate protocols (HTTP \rightleftharpoons gRPC, WebSocket upgrades).
 - **Payload Mapping & Shaping**
 - Fan-out or aggregate requests.
 - Rename fields, slim payloads, inject metadata.
 - **Performance Optimization**
 - Cache static and dynamic content at CDN edges.
 - Offload simple compute (geolocation, A/B logic) to edge workers.
 - **4.4 Ownership:** Platform team or shared between frontend/backend leads.
-

5. Backend Core Services (10 min)

- **5.1 Frontend-Integration Layer (BFF)**
 - Acts as the "true backend edge" for each client type (web, mobile, partner).
 - Tailors payloads, aggregates across domain services, and translates protocols (e.g. [gRPC](#) → REST).

- Enforces any client-specific caching, request shaping, or retry logic before hitting core services.
- **5.2 Business Domain Services**
 - Examples: Orders, Messages, Trips, Profiles.
 - Each service owns a cohesive set of entities and logic, exposing internal APIs for orchestration.
- **5.3 Service Mesh (Sidecar-Based Networking)**
 - Provides [mTLS](#), circuit-breaking, retries, and distributed tracing for all inter-service calls.
 - Enforces network-level policies ([ACLs](#), rate-limits) consistently, independent of service code.
 - Offloads common concerns so each service—plus the BFF layer—can remain lean.
- **5.4 Cross-Cutting Platform Services**
 - Shared infrastructure for auth, notifications, billing, metrics, experiments.
 - Some responsibilities (e.g. rate-limiting, token validation) may run at the edge or in BFF; heavier or async logic lives here.
- **5.5 Architecture Styles**
 - Monolith, modular monolith, or microservices: each can host BFFs and sidecars.
 - With microservices, you typically deploy BFFs and sidecars in the same cluster, but route traffic via the ingress/API gateway.
- **5.6 Communication Patterns**
 - Synchronous RPC (REST, [gRPC](#)) for BFF→Domain and Domain→Domain calls, all secured by the mesh.
 - Async Messaging (Pub/Sub, Queues) for background workflows, decoupled from client-facing BFFs.
- **5.7 Stateless vs Stateful Services**
 - BFFs and most domain services should be stateless to scale easily.
 - Stateful systems (sessions, game state) require sticky sessions or external state stores.
- **5.8 Infrastructure Patterns**
 - Sidecar proxies ([Envoy/Istio](#)) alongside every service pod, including BFF deployments.
 - Mesh control plane on its own tier, speaking to sidecars across all services.
- **5.9 Trade-Offs**
 - BFF adds per-client code but simplifies core services.
 - Service Mesh gives consistency and security but adds operational complexity.
- **5.10 Communication Style: Request vs. Event Driven**
 - Request-Driven (e.g., REST, [gRPC](#)): Synchronous RPC-style communication between services. Well-suited for real-time flows where immediate feedback is needed, such as user actions or transactional updates.

- Event-Driven (e.g., queues, Pub/Sub, event streams): Asynchronous messaging used for background jobs, system decoupling, and resilience under load. Often backed by [Apache Kafka](#), [Apache Pulsar](#), or cloud-native equivalents.
 - Common Patterns:
 - Core services often use request-based RPC, while async workflows (e.g., billing, notifications, ML pipelines) are handled via events.
 - Events are also used to power downstream consumers: audit logs, analytics, or materialized views.
 - Most systems combine both: sync for control paths, async for data and side effects.
 - **Reference Articles:**
 - [Request-Driven vs. Event-Driven Architecture Overview](#)
 - [Request-Driven \(RESTful\) vs Event-Driven in Microservices](#)
 - [Event-Driven Architecture Fundamentals and Common Pitfalls](#)
 - [Exploring Event-Driven Architecture: A Beginner's Guide for Cloud Native Developers](#)
- **5.11 Ownership**
 - BFFs often owned by the feature or frontend-platform team.
 - Core services by domain teams.
 - Mesh control plane by the platform/SRE team.
-

6. Data Storage & Retrieval (5 min)

- **6.1 Data Store Categories & Limitations**
 - **Relational (SQL):** [PostgreSQL](#), [MySQL](#)
 - Use: [ACID transactions](#), complex joins, ad-hoc queries.
 - Limitations: Challenging to scale horizontally (sharding/cluster required); heavy schema migrations.
 - **Key-Value Stores:** [Redis](#), [DynamoDB](#)
 - Use: Ultra-low-latency get/put for sessions, caches, user state.
 - Limitations: No rich query or relationship support; minimal indexing.
 - **Document Stores:** [MongoDB](#), [Couchbase](#)
 - Use: JSON-style documents, flexible/evolving schemas.
 - Limitations: Limited joins/transactions; risk of denormalization and data duplication.
 - **Column-Family Stores:** [Cassandra](#), [HBase](#)
 - Use: High write throughput, wide-column models, time-series data.
 - Limitations: Eventual consistency by default; operational complexity (compaction, repair).
 - **Graph Databases:** [Neo4j](#), [Amazon Neptune](#)
 - Use: Deep relationship traversals (social graphs, recommendations).
 - Limitations: Poor horizontal write scalability; specialized query language.
 - **Time-Series DBs:** [InfluxDB](#), [TimescaleDB](#)
 - Use: Efficient ingestion & querying of timestamped metrics/events.
 - Limitations: Needs retention/downsampling; storage bloat risk.
 - **Search & Analytics Engines:** [Elasticsearch](#), [Meilisearch](#), [Typesense](#)
 - Use: Full-text search, [faceted filtering](#), aggregations.

- Limitations: Eventual consistency; heavy index rebuilds; high resource use.
- 6.2 When to Choose Which Store
 - **SQL**: Transactions + joins + strong consistency.
 - **Key-Value**: Simple lookups + caching + session data.
 - **Document**: Semi-structured or rapidly evolving schemas.
 - **Graph**: Complex many-to-many relationships, deep traversals.
 - **Time-Series**: High-frequency timestamped writes, windowed queries.
 - **Search**: Text search, analytics, log exploration.
 - **Caution**: *Don't pick NoSQL merely to avoid schema design: poor data models create hidden debt.*
- 6.3 **CAP Theorem** & ACID Trade-Offs
 - **CAP Theorem** (under network partition)
 - CP: Consistency + Partition tolerance (sacrifice availability).
 - AP: Availability + Partition tolerance (sacrifice consistency).
 - CA: Consistency + Availability (no partition tolerance).

See also [this post](#) for a diagrammatic discussion of the three CAP theorem options.
 - **ACID Transactions**
 - Atomicity: All or nothing; if any part of the transaction fails, roll back.
 - Consistency: Valid state transitions; all data integrity constraints and business rules are maintained throughout the transaction.
 - Isolation: No interference among concurrent transactions. Changes made by one transaction are not visible to other transactions until the first transaction is committed.
 - Durability: Persistence through failures; changes are permanently stored in the database, even if the system fails or crashes.
- 6.4 Blob Storage for Media
 - **S3, GCS, Azure Blob**: object stores for large binaries or static assets (fronted by CDN).
- 6.5 Search Layer
 - **Elasticsearch, Meilisearch, Typesense**: specialized engines for text search and analytics; evaluate SaaS vs. self-hosted.
- 6.6 Ownership
 - Shared among application teams (schema design, queries), data engineering (ETL/pipelines), and platform/infrastructure teams (operations, scaling).

7. Scalability, Reliability & Performance (5 min)

When designing for scale, it's useful to distinguish between **vertical scaling** (making a single instance more powerful) and **horizontal scaling** (adding more instances). Most modern systems aim to scale horizontally for long-term flexibility, but both strategies have a role.

- 7.1 **Vertical Scaling (Scale-Up)**

- Add more CPU, RAM, or IOPS to a single node.
 - Quickest path to performance gains in early-stage systems.
 - Simple to implement, but limited by hardware ceilings and diminishing returns.
 - **7.2 Horizontal Scaling (Scale-Out)**
 - Add more nodes behind a load balancer (e.g., web servers, DB replicas).
 - Often requires stateless services, sharded databases, or eventual consistency models.
 - Enables fault tolerance and high availability but increases architectural complexity.
 - **7.3 Reliability Techniques**
 - Autoscaling triggers & constraints (CPU, memory, custom metrics like queue length).
 - Scale-in safeguards (to avoid thrashing during transient load spikes).
 - Health checks (readiness and liveness probes) are critical for detecting broken or overloaded instances and safely removing them from load balancers.
 - Auto-Healing mechanisms (e.g., in [Kubernetes](#) or with auto-scaling groups) restart unhealthy services automatically without human intervention.
 - Ensure retries are safe and idempotent for critical operations (e.g., payments, orders); use [idempotency keys](#) or deduplication when needed.
 - Chaos engineering (e.g., [Chaos Monkey](#)-style fault injection) to proactively validate fail-over paths.
See [Principles of Chaos Engineering](#) for an overview.
 - Multi-Region Deployment protects against entire region outages and reduces latency for global users, but adds complexity in data replication, failover routing, and consistency guarantees.
 - Defining SLIs/SLOs (latency P95, error budgets) as the foundation for any auto-scale or circuit-breaker policies.
 - **7.4 Performance Optimizations**
 - Caching: CDN, [Redis](#), application-layer caching.
 - [Backpressure and Resiliency](#): queuing, retry (with backoff), and circuit breakers
 - Query Tuning, Connection Pooling, Compression.
 - Load/stress testing (e.g. [Gatling](#), [JMeter](#)) to validate horizontal-scale behavior before prod rollouts.
 - **7.5. Ownership:** Typically shared between SRE, infrastructure teams, and performance-focused engineers within product teams.
-

8. Deployment, Observability, Security & Ownership (as time allows)

- **8.1 CI/CD & Infrastructure as Code**
 - Automate build, test, and deployment (e.g., [GitHub Actions](#), [CircleCI](#)).
 - Manage infra as code (e.g., [Terraform](#), [Pulumi](#)) and configuration as code ([Ansible](#), [Helm](#), [Packer](#)).
- **8.2 Deployment Strategies**
 - Blue/Green: instant rollback, doubled infra cost.

- Canary: phased rollout, relies on robust metrics and traffic control.
Safe deployments rely on robust pre-prod validation and test environments — see Section 9.9 ("Testing and Pre-Prod Environments") for guidance on staging, load testing, and end-to-end validation.

- **8.3 Observability & Alerting**

- Metrics, logs, and traces (e.g., [Prometheus](#), [Grafana](#), [OpenTelemetry](#)).
- Dashboards, [SLIs/SLOs](#), and error budgets.
- Alerting and on-call workflows (e.g., [PagerDuty](#), [Opsgenie](#)).

- **8.4 Security Controls**

- Authentication & authorization ([OAuth 2.0](#), [RBAC](#)).
- Secrets management (e.g., [HashiCorp Vault](#), [AWS Secrets Manager](#)).
- Network security (e.g., firewalls, mTLS).
- Encryption in transit ([TLS](#)) and at rest ([AES-256](#)).
- Supply chain integrity (image signing [Notary v2](#), vulnerability scanning [Trivy](#)).

- **8.5 Ownership**

- CI/CD & IaC: DevOps/Platform teams.
- Observability: SRE/Infrastructure teams.
- Security: Security team with feature-team integration.

9. Extensions, Trade-Offs, Bottlenecks & Evolution (as time allows)

- **9.1 Extensions & Integrations**

- **AI/ML Services**
 - Traditional models: recommendations, fraud detection, anomaly detection.
 - Generative AI: LLM-based summarization, chatbots, content generation.
- **Vector & Semantic Search**
 - Embedding pipelines, vector databases ([Pinecone](#), [Milvus](#)) for similarity search.
- **External APIs & SaaS Integrations**
 - Payment gateways, geolocation, identity providers, mapping, social logins.
- **Event-Driven Pipelines**
 - Streaming analytics, real-time notifications, data pipelines ([Apache Kafka](#), [Apache Pulsar](#)).
 - **Geographic Scaling:** multi-region failover, data residency, latency optimization.
 - **Privacy & Compliance:** [PII separation](#), encryption requirements, [GDPR/HIPAA](#) constraints.

- **9.2 Common Bottlenecks & Mitigations**

See Sections 7 ("Scalability, Reliability & Performance") and 6.1 ("Data Store Categories & Limitations") for patterns on write limits, cache pressure, queue backlogs, and latency mitigation.

- **9.3 Migration & Evolution Strategies**

- **Monolith to Microservices:** incremental extraction ([strangler-fig pattern](#)).

- **Modular Monolith:** adopt modules before splitting into services. ["monolith first" model](#)
 - Start with a well-modularised monolith, postpone the complexity premium of microservices, and let real usage teach you where the service boundaries should be before you carve them out.
- **9.4 Trade-Offs & Failure Modes**

Covered in Sections 6.3 ("CAP Theorem & ACID Trade-Offs"), 7.4 ("Performance Optimizations"), and 8.4 ("Security Controls").
- **9.5 Tenancy Models**
 - **Single-Tenant Systems:** Each customer has a dedicated instance of the application and its data. Offers strong isolation and simplified debugging but incurs higher operational cost.
 - **Multi-Tenant Systems:** Multiple customers share infrastructure, with logical data separation via tenant IDs. More cost-effective and scalable, but introduces complexity around data access, security boundaries, and performance isolation.
 - **Hybrid Models:** Mix of pooled and siloed tenants — e.g., shared app layer but isolated databases or compute for high-paying customers.
 - **Design Considerations:**
 - How is tenant identity propagated and enforced throughout the system? Apply fine-grained authorization (e.g., RBAC, per-object ACLs) where required.
 - What per-tenant quotas, rate limits, or billing hooks exist?
 - Can operational visibility be segmented by tenant?
 - Your architecture should account for tenant isolation, scaling patterns, and operational concerns from the beginning, even if you start single-tenant.
 - **Reference Articles:**
 - [Microsoft Learn: Tenancy Models](#)
 - [AWS Whitepaper: SaaS Tenant Isolation Strategies](#)
- **9.6 Platformization & Developer Experience**
 - **Onboarding & Scaffolding:** provide code and infra templates, guided tutorials, and [create-service](#) generators to minimize ramp-up time.
 - **Local Development Experience:** enable reproducible dev environments (containers, proxies), fast build and run loops, hot-reloading, and integrated debuggers.
 - **Developer Tooling & CLI:** offer a unified CLI or SDK for common tasks (service creation, infrastructure provisioning, feature flag toggling, logs/metrics retrieval).
 - **Documentation & Discoverability:** maintain searchable, versioned docs, interactive API catalogs, architecture diagrams, and example code to reduce cognitive load.
 - **CI/CD & Fast Feedback:** implement quick pipelines with pre-commit/lint checks, unit and integration tests, PR previews, and actionable build/test results to shift quality left.
 - **Observability Integration:** bake in structured logging, metrics, distributed tracing, and alerting defaults into new services, plus standardized dashboard and alert templates.
 - **Testing Infrastructure:** provide test data generators, mocking frameworks, contract testing, and local integration or ephemeral environments for reliable verification.
 - **Standards & Governance:** enforce code style, security policies, and compliance automatically via shared libraries, config-as-code, and pre-built CI/CD gates.

- **Self-Service Infrastructure:** enable developers to provision databases, queues, feature flags, and other resources through an internal portal or CLI without manual ops tickets.
- **Developer Metrics & Improvement Loops:** track key metrics (build/test times, deployment frequency, lead time, failure rates) and gather developer satisfaction feedback to drive iterative DX enhancements.
- **Culture & Community:** foster internal communities of practice, office hours, mentorship, and peer reviews to share knowledge and continuously evolve best practices.
- See generally: [Links to DevEx articles](#)

• 9.7 Organizational Impact & Team Structure

- **Conway's Law:** Recognize that system architecture often mirrors the organization's communication structure. Design choices can be constrained by, or deliberately influence, team boundaries.
See also: [Martin Fowler on Conway's Law](#).
- **Architecture & Teams:** Different architectures lend themselves to different team structures and require distinct skill sets.
 - **Microservices:** Often enable smaller, autonomous teams but require mature platform capabilities and operational practices. Consider the impact on inter-team communication and contracts.
See: [Team Topologies](#) for patterns.
 - **Monoliths:** May start simpler organizationally but can lead to development bottlenecks and complex dependencies as the system and team scale.
- **Skills Alignment:** Ensure the team possesses or can acquire the necessary skills (e.g., distributed systems, specific data stores, operational tooling) demanded by the chosen architecture.

• 9.8 Testing & Pre-Prod Environments

- Mature systems rely on multiple environments to validate changes safely before production.
- Common environments include:
 - **Staging:** Mirrors production as closely as possible for full end-to-end validation.
 - **Integration/Test:** Used to validate contract compatibility and shared services across teams.
 - **Load Testing:** Simulates peak traffic and failure conditions to verify scaling limits.
 - **Synthetic or Ephemeral Environments:** Temporary, isolated deployments (e.g., per-PR or feature branch) for faster iteration.
- **Testing Strategies:**
 - Unit tests for individual components.
 - Integration tests for service-to-service behavior.
 - End-to-end tests for full system workflows (often via UI automation or API smoke tests).
 - Contract tests (e.g., [PACT](#)) to detect breaking changes in upstream/downstream services.
- Good pre-prod hygiene includes test data seeding, schema evolution checks, and observability parity with prod.
- See: [Environment Overview](#)